

MARIOH: Multiplicity-Aware Hypergraph Reconstruction (Supplementary Document)

A. Negative Sampling

We generate hyperedge candidates for classifier training. Positive samples consist of true hyperedges taken directly from $\mathcal{H}^{(S)}$. Negative samples, on the other hand, are cliques of the same size as the true hyperedges but are not present in $\mathcal{H}^{(S)}$. These negative samples are obtained by sampling subcliques from larger cliques. Specifically, maximal cliques in the projected graph after filtering $\mathcal{G}'$ (as detailed in Sect. III-B) are grouped by size. For each true hyperedge size, we attempt to generate negative samples by randomly selecting subcliques of the same size from cliques larger than the true hyperedge size. If no larger cliques are available, we may not generate negative samples of that size, resulting in fewer negative samples than desired. We used a negative sampling ratio of 1 : 4, meaning that for each positive sample, we generated at most four negative samples.

Lemma 1 (Time Complexity of Negative Sampling). *Given graph $\mathcal{G}$, and N_p ground truth hyperedges with maximum size $s_{\max}$, the negative sampling algorithm generates $O(N_p)$ negative samples in total time*

$$O(C_{\text{total}} + N_p \cdot s_{\max}),$$

where C_{total} is the total number of maximal cliques in graph $\mathcal{G}$, and $s_{\max}$ is the maximum hyperedge size.

Proof. The negative sampling algorithm begins by finding all maximal cliques in the graph $\mathcal{G}$, which takes $O(C_{\text{total}})$ time [1], where C_{total} denotes the total number of maximal cliques. These cliques are then organized by their sizes using a hash map, requiring an additional $O(C_{\text{total}})$ time. For each ground truth hyperedge h of size s , the algorithm attempts to generate up to K (spec., 4) negative samples of the same size. It collects all maximal cliques of size at least s ; if no such cliques exist, it cannot generate negative samples for hyperedges of size s and proceeds to the next hyperedge.

When larger cliques are available, the algorithm makes up to $K \times 5$ attempts to generate negative samples for each h . In each attempt, it randomly selects a larger clique C and samples s nodes to form a candidate hyperedge h' . It then checks whether h' is neither in the ground truth set $\mathcal{H}^{(S)}$ nor already sampled; this check takes $O(1)$ time using efficient data structures like hash sets. If h' is valid, it is added to the set of negative samples. The attempt counter ensures that the algorithm does not run indefinitely when valid negative samples are scarce, particularly if many sampled sub-cliques are already in the ground truth or previously sampled.

Since K are constants, the time spent per hyperedge h is $O(s)$, where $s \leq s_{\max}$. Summing over all N_p ground truth hyperedges, the total time for negative sample generation is $O(N_p \cdot s_{\max})$. Therefore, the overall time complexity of the negative sampling algorithm is $O(C_{\text{total}} + N_p \cdot s_{\max})$. $\square$

B. Multiplicity-Aware Classifier

We use a 2-layered Multi-Layer Perceptron (MLP) to predict whether a given clique is a true hyperedge. The input to the classifier is the multiplicity-aware feature vector. The output layer uses a softmax activation to produce a probability distribution over two classes. The classifier is trained to minimize the binary cross-entropy loss using the predicted and ground-truth labels (1 for true hyperedges, 0 for non-hyperedges). Hidden layer dimensions and dropout rates are optimized depending on the datasets. A detailed hyper-parameter search space is described below.

Hyperparameter Search: We optimized the hyperparameters of the classifier using Bayesian optimization with the Optuna library [2]. Negative sampling was employed to generate the validation dataset from $\mathcal{G}^{(S)}$, and 5-fold cross-validation was performed to identify the hyperparameters that maximize AUROC. The search space included:

- **Batch Size:** {256, 512, 1024, 2048, 4096}
- **Number of Epochs:** [1, 50]
- **Learning Rate:** [10^{-5} , 10^{-2}]
- **Hidden Layer Dimensions:** {(128, 64), (128, 128), (256, 128), (256, 256)}
- **Dropout Rate:** [0.1, 0.5]

After 100 trials, the best hyperparameters achieving the highest mean AUROC during cross-validation were selected.

C. Baseline Methods for Hypergraph Reconstruction

• Overlapping Community Detection

- **CFINDER** [3] identifies communities by finding overlapping k -cliques. The optimal k is selected using a Jaccard-based metric on the training set.
- **DEMON** [4] identifies overlapping communities in a graph by iteratively aggregating local communities discovered around each node. It uses a threshold parameter ϵ to merge similar communities and a minimum community size to filter out small communities.

• Clique Decomposition

- **CLIQUE COVERING** [5] covers all edges in a graph by identifying cliques that include uncovered edges. It iteratively selects an uncovered edge, expands it into a

TABLE I: Execution time comparison. Average runtime (in seconds) across datasets, reported with standard deviations over five runs. "OOT" indicates "out of time" (exceeding 24 hours), and "OOM" represents "out of memory" (system limit: 384 GB).

Dataset	CFINDER	DEMON	CLIQUE COVERING	MAX CLIQUE	BAYESIAN-MDL	SHYRE-UNSUP	SHYRE-Motif	SHYRE-Count	MARIOH
Enron	0.23 ± 0.00	0.15 ± 0.00	0.01 ± 0.00	0.03 ± 0.00	0.46 ± 0.00	1.55 ± 0.29	22.89 ± 0.69	2.18 ± 0.30	5.15 ± 0.54
PSchool	277.44 ± 2.86	0.83 ± 0.00	0.06 ± 0.00	0.67 ± 0.02	9.19 ± 0.09	50.63 ± 0.44	OOT	16.84 ± 1.04	27.98 ± 0.89
H.School	3.20 ± 0.02	0.48 ± 0.01	0.03 ± 0.00	0.13 ± 0.00	1.74 ± 0.01	10.49 ± 0.13	1281.12 ± 23.79	4.78 ± 0.65	5.98 ± 0.06
Crime	0.02 ± 0.00	0.15 ± 0.02	0.01 ± 0.00	0.01 ± 0.00	0.39 ± 0.00	0.42 ± 0.00	0.27 ± 0.04	0.20 ± 0.02	0.79 ± 0.02
Hosts	0.56 ± 0.00	2.30 ± 0.02	0.01 ± 0.00	0.06 ± 0.00	2.45 ± 0.01	3.91 ± 0.01	336.64 ± 5.33	1.70 ± 0.37	4.17 ± 0.16
Directors	0.02 ± 0.00	0.12 ± 0.00	0.01 ± 0.00	0.01 ± 0.00	0.38 ± 0.00	0.10 ± 0.02	0.11 ± 0.01	0.14 ± 0.02	0.86 ± 0.18
Foursquare	8.56 ± 0.10	9.70 ± 0.11	0.16 ± 0.00	0.65 ± 0.00	28.79 ± 0.60	22.87 ± 0.17	8422.65 ± 259.82	3.72 ± 0.19	8.19 ± 0.22
DBLP	29.18 ± 0.57	18143.66 ± 579.85	8.46 ± 0.18	11.99 ± 0.67	245.16 ± 0.61	OOT	3046.27 ± 104.93	2677.27 ± 173.34	658.46 ± 4.97
Eu	16363.26 ± 2186.62	5.34 ± 1.34	0.29 ± 0.04	12.20 ± 6.15	36.77 ± 13.47	1925.37 ± 200.26	OOM	246.23 ± 1.86	1530.42 ± 218.24
TopCS	4.33 ± 0.07	270.03 ± 1.54	1.17 ± 0.01	1.54 ± 0.02	42.01 ± 0.33	5889.73 ± 87.56	1068.01 ± 14.62	101.36 ± 2.77	84.30 ± 3.89

maximal clique using neighboring nodes, and removes edges covered by the clique until all edges are included in at least one clique.

- **MAX CLIQUE** [1] considers all maximal cliques in the graph as hyperedges.

• Hypergraph Reconstruction

- **BAYESIAN-MDL** [6] reconstructs hypergraphs using a Bayesian generative model, emphasizing the principle of parsimony to infer the simplest hypergraph that explains the observed pairwise data. It calculates the posterior probability of a hypergraph and estimates structures through Markov Chain Monte Carlo sampling.
- **SHYRE** [7] reconstructs hypergraphs by sampling cliques from the projected graph based on a statistical distribution $\rho(n, k)$, which captures patterns of hyperedges within maximal cliques. A classifier is trained to distinguish between hyperedges and non-hyperedges. SHYRE-Count uses basic structural features, while SHYRE-Motif incorporates subgraph patterns such as triangle and square motifs as additional features for classification.
- **SHYRE-UNSUP** [7] reconstructs hypergraphs by iteratively selecting maximal cliques as candidates for hyperedges. The cliques are ranked based on their size and average edge multiplicity, preferring larger cliques with lower average edge multiplicities. Once selected, a clique is converted into a hyperedge, and the corresponding edge multiplicities in the projected graph are reduced until all edge multiplicities are satisfied.

D. Bidirectional Search Parameters

The initial classification threshold θ_{init} and the negative prediction processing ratio r were determined as the combination that maximized reconstruction accuracy on the source projected graph $G^{(S)}$ and its corresponding source hypergraph $H^{(S)}$. Specifically, hyperparameters are explored from the ranges: $\theta_t = [0.5, 0.6, \dots, 1.0]$, $r = [5, 10, \dots, 100]$. For all experiments, except those conducted for the parameter sensitivity analysis (Fig. 3 in the main paper), MARIOH was configured with $\alpha = 1/20$.

E. Structural Properties of Hypergraphs

We analyzed the structural characteristics of hypergraphs by quantifying and comparing the following metrics across the ground truth hypergraph ($\mathcal{H}^{(T)}$) and the reconstructed hypergraph ($\hat{\mathcal{H}}^{(T)}$):

- **Number of Nodes and Hyperedges:** The total count of unique nodes and hyperedges in the hypergraph was computed to capture its overall size and complexity. This metric provides fundamental insights into the scale and sparsity of the hypergraph.
- **Mean Hyperedge Sizes:** We calculated the mean (μ) of the hyperedge sizes to measure their central tendency and variability. These metrics reflect the heterogeneity in the hypergraph's structural composition.
- **Degree Distributions (Single, Pairwise, Triple):** The degree distributions of nodes, node pairs, and node triples within hyperedges were analyzed to characterize their occurrence frequencies. This metric highlights how nodes participate in single hyperedges, pairwise connections, and higher-order interactions [8]–[10].
- **Density and Overlapness:** Hypergraph density was defined as the ratio of hyperedges to nodes [11], while overlapness was measured as the ratio of the total hyperedge size to the number of nodes [10]. These metrics provide a macroscopic view of the connectivity and overlap of hyperedges within the hypergraph.
- **Simplicial Closure Ratio:** This metric measures the proportion of closed triangles formed by three mutually connected nodes within the hypergraph. A higher closure ratio indicates a greater tendency for nodes to form tightly connected substructures, which is an essential characteristic of higher-order networks [12].
- **Hyperedge Homogeneity:** To evaluate the internal consistency of hyperedges, we quantified the average pairwise connection frequency within each hyperedge [10]. This metric captures the uniformity of relationships among hyperedge members.
- **Singular Value Distribution:** We computed the singular values of the adjacency matrix derived from the projected graph representation of the hypergraph. The distribution of singular values provides insights into the structural connectivity and dimensionality of the graph [9], [13].

F. Runtime Analysis for Individual Datasets

The execution times presented in Table I highlight significant differences in runtime efficiency among the methods across various datasets. MARIOH consistently demonstrates competitive performance, with moderate execution times across all datasets. SHYRE-Count generally exhibits slightly faster runtimes than MARIOH, while SHYRE-Motif and SHYRE-UNSUP show considerably higher runtimes, often reaching Out-of-Time (OOT) or Out-of-Memory (OOM) limits

for larger datasets like DBLP and EU. BAYESIAN-MDL achieves the fastest execution times among the hypergraph reconstruction methods, followed closely by MARIOH.

Clique decomposition methods (MAX CLIQUE and CLIQUE COVERING) are the fastest overall but at the cost of significantly lower reconstruction accuracy (as shown in Table II). Community-based methods like DEMON and CFINDER exhibit much higher runtimes for large datasets, with CFINDER being particularly slow on datasets such as EU and TopCS. These results highlight MARIOH’s balance between runtime efficiency and reconstruction accuracy, making it a reliable option for hypergraph analysis.

REFERENCES

- [1] E. Tomita, A. Tanaka, and H. Takahashi, “The worst-case time complexity for generating all maximal cliques and computational experiments,” *Theoretical computer science*, vol. 363, no. 1, pp. 28–42, 2006.
- [2] T. Akiba, S. Sano, T. Yanase, T. Ohta, and M. Koyama, “Optuna: A next-generation hyperparameter optimization framework,” in *KDD*, 2019.
- [3] G. Palla, I. Derényi, I. Farkas, and T. Vicsek, “Uncovering the overlapping community structure of complex networks in nature and society,” *nature*, vol. 435, no. 7043, pp. 814–818, 2005.
- [4] M. Coscia, G. Rossetti, F. Giannotti, and D. Pedreschi, “Demon: a local-first discovery method for overlapping communities,” in *KDD*, 2012.
- [5] A. Conte, R. Grossi, and A. Marino, “Clique covering of large real-world networks,” in *SAC*, 2016.
- [6] J.-G. Young, G. Petri, and T. P. Peixoto, “Hypergraph reconstruction from network data,” *Communications Physics*, vol. 4, no. 1, p. 135, 2021.
- [7] Y. Wang and J. Kleinberg, “From graphs to hypergraphs: Hypergraph projection and its reconstruction,” in *ICLR*, 2024.
- [8] M. T. Do, S.-e. Yoon, B. Hooi, and K. Shin, “Structural patterns and generative models of real-world hypergraphs,” in *KDD*, 2020.
- [9] Y. Kook, J. Ko, and K. Shin, “Evolution of real-world hypergraphs: Patterns and models without oracles,” in *ICDM*, 2020.
- [10] G. Lee, M. Choe, and K. Shin, “How do hyperedges overlap in real-world hypergraphs?—patterns, measures, and generators,” in *WWW*, 2021.
- [11] S. Hu, X. Wu, and T. H. Chan, “Maintaining densest subsets efficiently in evolving hypergraphs,” in *CIKM*, 2017.
- [12] A. R. Benson, R. Abebe, M. T. Schaub, A. Jadbabaie, and J. Kleinberg, “Simplicial closure and higher-order link prediction,” *Proceedings of the National Academy of Sciences*, vol. 115, no. 48, pp. E11 221–E11 230, 2018.
- [13] M. E. Wall, A. Rechtsteiner, and L. M. Rocha, “Singular value decomposition and principal component analysis,” in *A practical approach to microarray data analysis*. Springer, 2003, pp. 91–109.